



CODING HORROR
programming and human factors



Transcreva áudio em texto de forma rápida e precisa com a plataforma de IA do Azure.

ADS VIA CARBON

RESOURCES


About Me

discourse.org

stackexchange.com

[Learn Markdown](#)

[Recommended Reading](#)

 [Subscribe in a reader](#)

 [Subscribe via email](#)

Coding Horror has been continuously published since 2004

Copyright Jeff Atwood © 2022

Logo image © 1993 Steven C. McConnell

Proudly published with  Ghost

11 Oct 2007

A Visual Explanation of SQL Joins

I thought Ligaya Turmelle's [post on SQL joins](#) was a great primer for novice developers. Since SQL joins *appear* to be set-based, the use of [Venn diagrams](#) to explain them seems, at first blush, to be a natural fit. However, like the commenters to her post, I found that the Venn diagrams didn't quite match the [SQL join syntax](#) reality in my testing.


I love the concept, though, so let's see if we can make it work. Assume we have the following two tables. **Table A** is on the left, and **Table B** is on the right. We'll populate them with four records each.

| id | name | id | name |
|----|-----------|----|-------------|
| -- | ---- | -- | ---- |
| 1 | Pirate | 1 | Rutabaga |
| 2 | Monkey | 2 | Pirate |
| 3 | Ninja | 3 | Darth Vader |
| 4 | Spaghetti | 4 | Ninja |

Let's join these tables by the name field in a few different ways and see if we can get a conceptual match to those nifty Venn diagrams.

```
SELECT * FROM TableA
INNER JOIN TableB
ON TableA.name = TableB.name
```

| id | name | id | name |
|----|--------|----|--------|
| 1 | Pirate | 2 | Pirate |
| 3 | Ninja | 4 | Ninja |



Inner join produces only the set of records that match in both Table A and Table B.

```
SELECT * FROM TableA
FULL OUTER JOIN TableB
ON TableA.name = TableB.name
```

| id | name | id | name |
|----|------|----|------|
|----|------|----|------|

| | | | |
|---|--------|---|--------|
| 1 | Pirate | 2 | Pirate |
|---|--------|---|--------|

| | | | |
|---|--------|------|------|
| 2 | Monkey | null | null |
|---|--------|------|------|

| | | | |
|---|-------|---|-------|
| 3 | Ninja | 4 | Ninja |
|---|-------|---|-------|

| | | | |
|---|-----------|------|------|
| 4 | Spaghetti | null | null |
|---|-----------|------|------|

| | | | |
|------|------|---|----------|
| null | null | 1 | Rutabaga |
|------|------|---|----------|

| | | | |
|------|------|---|-------------|
| null | null | 3 | Darth Vader |
|------|------|---|-------------|



Table A

Full outer join produces the set of all records in Table A and Table B, with matching records from both sides where available. If there is no match, the missing side will contain null.

```
SELECT * FROM TableA
LEFT OUTER JOIN TableB
ON TableA.name = TableB.name
```

| id | name | id | name |
|----|-----------|------|--------|
| -- | ---- | -- | ---- |
| 1 | Pirate | 2 | Pirate |
| 2 | Monkey | null | null |
| 3 | Ninja | 4 | Ninja |
| 4 | Spaghetti | null | null |



Table A

Left outer join produces a complete set of records from Table A, with the matching records (where available) in Table B. If there is no match, the right side will contain null.

```
SELECT * FROM TableA
LEFT OUTER JOIN TableB
ON TableA.name = TableB.name
WHERE TableB.id IS null
```

| id | name | id | name |
|----|------|----|------|
|----|------|----|------|

| | | | |
|---|--------|------|------|
| 2 | Monkey | null | null |
|---|--------|------|------|

| | | | |
|---|-----------|------|------|
| 4 | Spaghetti | null | null |
|---|-----------|------|------|



Table A

To produce the set of records only in Table A, but not in Table B, we perform the same left outer join, then **exclude the records we don't want from the right side via a where clause.**

```
SELECT * FROM TableA
FULL OUTER JOIN TableB
ON TableA.name = TableB.name
WHERE TableA.id IS null
OR TableB.id IS null
```

| id | name | id | name |
|----|------|----|------|
|----|------|----|------|

| | | | |
|---|--------|------|------|
| 2 | Monkey | null | null |
|---|--------|------|------|

| | | | |
|---|-----------|------|------|
| 4 | Spaghetti | null | null |
|---|-----------|------|------|

| | | | |
|------|------|---|----------|
| null | null | 1 | Rutabaga |
|------|------|---|----------|

| | | | |
|------|------|---|-------------|
| null | null | 3 | Darth Vader |
|------|------|---|-------------|



Table A

To produce the set of records unique to Table A and Table B, we perform the same full outer join, then **exclude the**

records we don't want from both sides via a where clause.

There's also a cartesian product or **cross join**, which as far as I can tell, can't be expressed as a Venn diagram:

```
SELECT * FROM TableA  
CROSS JOIN TableB
```

This joins "everything to everything", resulting in $4 \times 4 = 16$ rows, far more than we had in the original sets. If you do the math, you can see why this is a *very* dangerous join to run against large tables.

NEXT

Mouse Ballistics

PREVIOUS

A Lesson in Control Simplicity

Written by Jeff Atwood

Indoor enthusiast. Co-founder of Stack Overflow and Discourse. Disclaimer: I have no idea what I'm talking about. Find me here: <http://twitter.com/codinghorror>